

PART 1

Define a hypothetical AI problem.

To define a hypothetical AI problem, I have considered "**Predicting Student Dropout Rates**" as my focus area. This problem is significant in the educational sector, where understanding and mitigating dropout rates can lead to improved student retention and institutional performance.

List 3 objectives and 2 stakeholders.

- ② **Identify Risk Factors:** Develop a model that analyzes various academic, demographic, and financial factors to identify students at high risk of dropping out. This includes examining grades, attendance, socioeconomic status, and engagement levels.
- ② **Implement Early Intervention Strategies:** Use the predictive model to inform educational institutions about at-risk students, enabling timely interventions such as academic counseling, financial aid, or personalized support programs.
- ② **Evaluate Model Effectiveness:** Continuously assess the accuracy and effectiveness of the predictive model by comparing predicted dropout rates with actual outcomes, refining the model based on feedback and new data.

Stakeholders

- **Educational Institutions:** Universities and colleges that aim to improve student retention rates and reduce financial losses associated with high dropout rates.
- **Students:** The primary beneficiaries of the predictive model, as it aims to provide them with the support needed to succeed academically and remain enrolled.

Propose 1 Key Performance Indicator (KPI) to measure success.

- **Dropout Prediction Accuracy:** Measure the accuracy of the predictive model by calculating the percentage of correctly identified at-risk students compared to actual dropout rates. This KPI will help gauge the model's effectiveness in predicting student dropout and inform necessary adjustments to improve its predictive capabilities.

By focusing on these objectives, stakeholders, and KPIs, the initiative to predict student dropout rates can be effectively structured to enhance educational outcomes and support student success.

Identify 2 data sources for your problem.

② **Academic Management System (AMS)**: Collects demographic, enrollment, and academic performance data.

② **Learning Management System (LMS)**: Tracks student engagement and learning behaviors, such as assignment submissions and forum participation.

Explain 1 potential bias in the data.

- **Selection Bias**: If the dataset over-represents certain demographics, the model may not generalize well, potentially underestimating dropout risks for underrepresented groups.

Outline 3 preprocessing steps.

② **Data Cleaning**: Remove duplicates and handle missing values to ensure data accuracy.

② **Feature Engineering**: Create new features, like a "student engagement score," to enhance predictive power.

② **Normalization and Encoding**: Normalize numerical features and encode categorical variables for compatibility with machine learning algorithms.

Choose a model (e.g., Random Forest, Neural Network) and justify your choice.

For predicting student dropout rates, I recommend using the **Random Forest** model due to its high accuracy and robustness against overfitting. It effectively combines multiple decision trees to enhance prediction stability and accuracy, making it suitable for educational data analysis.

Describe how you would split data into training/validation/test sets.

I would split the dataset as follows:

- **Training Set (70%)**: For training the model.
- **Validation Set (15%)**: For tuning hyperparameters.
- **Test Set (15%)**: For evaluating model performance on unseen data

Name **2 hyperparameters** you would tune and why.

② **Number of Trees (n-estimators)**: This controls the number of trees in the forest, impacting accuracy and computation time.

③ **Maximum Depth (max-depth)**: This limit how deep each tree can grow, helping to prevent overfitting while maintaining model performance.

Select **2 evaluation metrics** and explain their relevance.

- **F1 Score**: This metric balance precision and recall, making it essential for dropout prediction where class imbalance exists. It ensures that both the identification of at-risk students and the minimization of false positives are effectively managed.
- **Precision-Recall AUC**: This metric is particularly useful in imbalanced datasets, focusing on the model's ability to accurately identify positive instances (dropouts) while reducing false positives, which is critical in educational interventions.

What is **concept drift**? How would you monitor it post-deployment?

Concept drift refers to changes in the relationship between input features and the target variable over time. It can lead to decreased model accuracy as the patterns learned during training become outdated. To monitor concept drift post-deployment, one can track model performance metrics and average confidence scores over time, looking for significant changes that indicate drift.

Describe **1 technical challenge** during deployment (e.g., scalability).

A major challenge is **scalability**. As the volume of data increases, the model must efficiently process larger datasets without sacrificing performance. This often requires robust infrastructure and potentially distributed computing solutions to maintain real-time accuracy across diverse student populations.

Part 2: Case Study Application.

Problem Scope

Definition: Develop an AI system to predict patient readmission risk within 30 days post-discharge.

Objectives:

- Identify high-risk patients.
- Enable timely interventions to reduce readmissions.
- Improve patient outcomes and reduce costs.

Stakeholders:

- Patients: Benefit from personalized care.

- Healthcare Providers: Use insights for better decision-making.
- Hospital Administration: Focus on cost reduction and quality improvement.

Data Strategy

Data Sources:

- EHRs: Clinical data including diagnoses and treatment history.
- Demographics: Age, gender, socioeconomic status.
- Wearable Devices: Continuous health monitoring data.

Ethical Concerns:

- Patient Privacy: Protecting sensitive health information.
- Bias: Ensuring fair treatment across diverse patient populations.

Preprocessing Pipeline:

- Data Cleaning: Remove duplicates and correct errors.
- Feature Engineering: Create new features from existing data, normalize and encode variables.
- Data Splitting: Divide into training, validation, and test sets.

Model Development

Model Selection:

- Random Forest: Chosen for its robustness and ability to handle various data types.

Confusion Matrix and Metrics (Hypothetical Data):

- TP: 80, TN: 150, FP: 20, FN: 50.

Confusion Matrix: Predicted Positive Predicted Negative

Actual Positive	80	50
-----------------	----	----

Actual Negative	20	150
-----------------	----	-----

Predicted Positive	Predicted Negative
--------------------	--------------------

Actual Positive	80	50
-----------------	----	----

Actual Negative	20	150
-----------------	----	-----

Calculations:

* Precision: $80/(20+80)=0.80$

* Recall: $80/(80 + 50)= 0.61$

Integration Steps:

1. Integrate the model into hospital systems.

2. Train staff on using the AI tool.
3. Monitor performance and adjust as needed.

Compliance:

- Ensure adherence to HIPAA through data encryption and access controls.

Overfitting Solution:

- Cross-Validation: Use k-fold cross-validation to ensure model generalization.

Part 3: Critical Thinking.

Impact of Biased Training Data on Patient Outcomes.

Biased training data can lead to inaccurate predictions and unequal treatment recommendations, particularly affecting underrepresented groups. This can result in misdiagnoses and exacerbate health disparities, as certain populations may not receive appropriate care based on their unique health profiles.

Strategy to Mitigate Bias.

To mitigate bias, ensure diverse and representative datasets during training. This involves actively collecting data from various demographic groups and implementing rigorous testing protocols to identify and address biases.

Trade-offs Between Model Interpretability and Accuracy.

There is a trade-off between model interpretability and accuracy in healthcare. Complex models may achieve higher accuracy but lack transparency, making it difficult for clinicians to trust their recommendations. Simpler, more interpretable models may sacrifice some accuracy but enhance user confidence and understanding.

Impact of Limited Computational Resources on Model Choice.

Limited computational resources may necessitate the use of simpler models that require less processing power. While these models may be less accurate, they are easier to interpret and integrate into clinical workflows, ensuring that healthcare providers can effectively utilize them without overwhelming existing infrastructure.

Part 4: Reflection & Workflow Diagram

Reflection

Most Challenging Part of the Workflow

The most challenging aspect of the AI development workflow is **data collection and preparation**. This phase is critical because poor data quality can lead to biased models and inaccurate predictions.

Improvement with More Time/Resources

With additional time and resources, I would enhance data collection by implementing **automated tools** for data acquisition and **advanced cleaning techniques** to ensure high-quality, representative datasets.

Diagram: AI Development Workflow

[Problem Definition]



[Data Collection]



[Data Preparation]



[Model Design]



[Model Training]



[Model Evaluation]



[Deployment]



[Monitoring and Maintenance]